

A Lightweight Serverless Execution Platform

Bachelor Project Prototype Documentation

A practical report on the design, implementation, deployment, evaluation, and future work of a Docker-based serverless platform.

Project: Lightweight Serverless Platform
Prototype name: Ephemeral Mist / Fleeting Circus deployment
Main technologies: Django, Redis Streams, PostgreSQL, Docker, React, Grafana
Deployment domain: `fleetingcircus.ir`
Date: September 2026

Contents

Preface	1
1 Platform Overview	2
1.1 What the platform is	2
1.2 Current user-facing capabilities	2
1.3 Current deployment shape	3
2 Using the Platform	4
2.1 Mental model for users	4
2.2 Source and dependency specification	4
2.3 Access modes	5
2.4 Asynchronous invocation	5
2.5 Synchronous invocation	5
2.6 Downloading invocation artifacts	6
3 Architecture	7
3.1 Service responsibilities	7
3.2 Data stores	7
3.2.1 PostgreSQL	7
3.2.2 Redis	8
3.2.3 Object storage	8
3.3 Why a separate userservice exists	8
4 Job Model and Orchestration	9
4.1 Why orchestration became necessary	9
4.2 Coordination versions	9
4.3 V2 job state	9
4.4 Why Redis Streams	10
4.5 Backoff and dead-letter behavior	10
5 Build Flow	11
5.1 Build overview	11
5.2 Source bundle	11
5.3 Builder output	11
5.4 Build image lifecycle	12
5.5 Build cancellation and retry	12

6	Invocation Flow in Detail	13
6.1	Asynchronous invocation, step by step	13
6.2	Invocation execution diagram	14
6.3	Finalization and projection diagram	14
6.4	Input sandbox	14
6.5	Output sandbox	15
6.6	Why staged artifacts exist	18
6.7	Simple and advanced response modes	18
7	The Worker and Function Containers	20
7.1	Worker responsibilities	20
7.2	Worker heartbeat	20
7.3	Container create versus start	21
7.4	Warm containers	21
7.5	Warm container limitations	21
8	Scheduling Policy	24
8.1	What the orchestrator actually owns	24
8.2	Job creation and ready set	24
8.3	Worker heartbeats as scheduling input	25
8.4	Separate build and invocation streams	25
8.5	Invocation placement policy	26
8.6	Guarded sticky fallback	27
8.7	Build placement policy	28
8.8	Atomic dispatch	28
8.9	Claim and lease protocol	29
8.10	Completion, finalization, and ACK	30
8.11	Recovery and dead letters	31
8.12	What can go wrong	31
9	Security and Access Control	32
9.1	Authentication	32
9.2	Authorization	32
9.3	Invocation tokens	32
9.4	Current hardening gaps	32
10	Artifacts, Logs, and Retention	33
10.1	Artifact categories	33
10.2	Log storage	33
10.3	Object storage	33
11	Reliability Design	34
11.1	Fencing tokens	34
11.2	Leases and recovery	34
11.3	Safe completion protocol	34
11.4	Graceful shutdown	35

12 Observability	36
12.1 Stack	36
12.2 Current observability limitation	36
12.3 Interpreting empty panels	36
13 Deployment	37
13.1 Local deployment	37
13.2 Distributed prototype deployment	37
13.3 Private network	37
13.4 Registry behavior	37
14 Benchmarks and Performance	38
14.1 Benchmark terminology	38
14.2 Warm-container benchmark	38
14.3 Distributed worker concurrency benchmark	39
14.4 Invocation latency profiling before warm-container tuning	40
14.5 Where time is currently spent	41
15 Things Tried, Deprecated, and Changed	42
15.1 Backend-owned coordination	42
15.2 Redis Lists	42
15.3 Docker export/copy for outputs	43
15.4 Keeping containers alive with sleep	43
15.5 Sticky routing alone	44
15.6 Single-service authentication	44
15.7 Direct runner upload tradeoff	44
15.8 Warm containers before resident runners	44
16 Current Bottlenecks and Improvement Plan	46
16.1 Most important bottlenecks	46
16.2 Highest-value improvements	47
16.3 Research direction	48
16.4 Dependency and wheel strategy	48
17 Reading the Repository	49
17.1 Recommended reading order	49
17.2 Backend code	49
17.3 Worker code	50
17.4 Scheduler and orchestrator code	50
17.5 Frontend code	51
18 Conclusion	52
A Selected API Reference	53
B Example Function With Output File	54
C Glossary	55

Preface

This document describes the current prototype of a lightweight serverless execution platform. It is written as a project guide rather than as a short marketing overview. The goal is that a new developer, examiner, or future maintainer can read this report together with the repository and understand what the system does, why it is built this way, what has already been tried, where the risks are, and where the next engineering work should go.

The prototype is intentionally not presented as a production cloud platform. It is a working research and engineering prototype: users can create Python functions, upload source code, build Docker images, invoke those functions, send JSON and files, receive results and artifacts, use invocation tokens, observe the system through Grafana, and run the stack locally or across a small distributed deployment. The most interesting part is not only that the path works, but that the project has gradually moved from a simple backend-driven queue to a more explicit orchestration design using Redis Streams, worker heartbeats, fenced attempts, staged artifacts, and warm containers.

The document is organized from the outside inward. It first explains what the platform is and what a user can do with it. It then describes the architecture, the job flows, the build and invocation pipeline, the orchestrator, workers, containers, artifact handling, reliability mechanisms, deployment, benchmarks, known bottlenecks, and repository layout.

1 Platform Overview

1.1 What the platform is

The platform is a small serverless execution environment for Python functions. A user writes a handler function, declares the shape of its inputs and outputs, builds it into an image, and invokes it through platform APIs. The caller does not run Docker directly and does not need to know where the worker machines are. The platform accepts the request, schedules the job, executes the function in an isolated container, and exposes the result through API endpoints.

The word “serverless” in this project does not mean that servers do not exist. It means that the user works with functions and requests instead of managing machines. The prototype still uses servers, Docker, Redis, PostgreSQL, and a registry, but those are platform concerns rather than user concerns.

1.2 Current user-facing capabilities

The current prototype supports the following user-facing operations:

- Register and log in through a separate userservice.
- Create and edit functions.
- Paste Python handler code and `requirements.txt` dependencies in the frontend.
- Declare whether a function is private, token-protected, or public.
- Declare allowed input files and expected output files.
- Build a function into a Docker image.
- View build status and build failures.
- Invoke functions asynchronously.
- Invoke eligible JSON-only functions synchronously.
- Send JSON events and, for asynchronous invocations, file inputs.
- Read final JSON results, stdout, stderr, exit status, and timing.
- List and download declared output files.

- Download an invocation ZIP containing the invocation reference, inputs, outputs, stdout, stderr, and result metadata.
- Create, rotate, revoke, expire, and list invocation tokens.
- Observe the deployed platform through Grafana, Prometheus, Loki, and exporters.

The main product idea is that the function owner controls who can call the function. JWT authentication is used for account and function management. Invocation tokens are separate secrets used by non-owner callers. This separation is important because a function owner may want to give one client access to a single function without giving that client access to their account.

1.3 Current deployment shape

The project supports two deployment modes. The local mode uses one Docker Compose stack and is useful for development and repeatable tests. The first distributed mode places the control plane on one machine and runs workers on separate machines over a private network.




Figure 1.1: Current distributed prototype topology.

2 Using the Platform

2.1 Mental model for users

A user works with four main concepts:

Function	A named Python unit owned by a user.
Source	The function code, dependencies, handler setting, and contract.
Build	A Docker image produced from the source.
Invocation	One request to run a built function with a specific event and optional files.

The function source must define a handler. A simple handler looks like this:

Listing 2.1: Minimal Python handler.

```
1 def main(event, context):
2     name = event.get("name", "world")
3     return {"message": f"Hello, {name}!"}
```

The handler receives two arguments. The `event` argument is the JSON payload supplied by the caller. The `context` argument contains platform metadata such as the request identifier and paths for input and output files.

2.2 Source and dependency specification

The frontend lets users paste the main handler code and `requirements.txt`. Internally the platform packages those files into a ZIP source bundle. The worker turns that bundle into a Docker image. The image build uses the declared runtime, currently Python 3.13 in the common path.

Listing 2.2: Example requirements.txt.

```
1 requests==2.32.3
2 numpy==2.2.0
```

Dependency builds are one of the areas where the prototype is useful but not yet complete. Pure Python packages are straightforward. Packages with compiled wheels or system-level dependencies may require extra build image support. Future work should include runtime base image variants, build templates, or a controlled mechanism for system packages.

2.3 Access modes

Each function has an invocation access mode:

- **private**: only the owner or an admin can invoke with a JWT.
- **token**: callers invoke with **X-Function-Token**.
- **public**: callers can invoke without JWT or function token.

Function tokens are intentionally not JWTs. A JWT proves account identity. A function invocation token proves permission to call one function. This makes it possible to rotate or revoke function access without changing account login tokens.

2.4 Asynchronous invocation

Asynchronous invocation is the main durable execution path. It supports file inputs, declared output files, polling, and ZIP download. The first response returns an invocation identifier and a read token. The caller polls the invocation endpoint until the invocation reaches a terminal state.

Listing 2.3: Asynchronous invocation example.

```
1 curl -X POST https://fleetingcircus.ir/api/functions/42/invoke/ \
2   -H "Authorization: Bearer <jwt>" \
3   -H "Content-Type: application/json" \
4   -d '{"event":{"name":"Ilya","numbers":[1,2,3]}}'
```

A simple response mode is the default. It keeps the API response focused on what most users want: status, returned result, and output paths. The advanced mode includes more platform metadata for debugging and frontend diagnostics.

2.5 Synchronous invocation

Synchronous invocation is a convenience path for functions that are expected to finish quickly and do not require file input or declared output files. It waits for the result and returns it in the original HTTP response if the function finishes within the synchronous wait budget. The function timeout is separate: it counts only actual function execution time, not network delay, queueing delay, or scheduling delay.

Listing 2.4: Synchronous invocation example.

```
1 curl -X POST https://fleetingcircus.ir/api/functions/42/invoke-sync/ \
2   -H "Authorization: Bearer <jwt>" \
3   -H "Content-Type: application/json" \
4   -d '{"event":{"x":10,"y":20},"response_mode":"simple"}'
```

2.6 Downloading invocation artifacts

For asynchronous invocations, the caller can download an invocation ZIP while the invocation is still retained. The retention rule for the prototype is that invocations, their inputs, outputs, and logs expire together after seven days. Expired invocations return HTTP 404 as if they never existed.

The ZIP contains a reference to the function and version, the original input files, declared output files, stdout, stderr, and result metadata. Source code is not included in invocation downloads.

3 Architecture

3.1 Service responsibilities

The architecture is deliberately split into services even though this is still a prototype. The separation reflects the direction of the system:

Service	Responsibility
<code>userservice</code>	Owns account data, password handling, JWT issuing, refresh, role claims, and public key/JWKS discovery.
<code>backend</code>	Owns the product API: functions, source bundles, builds, invocations, tokens, artifacts, downloads, OpenAPI, and ownership checks.
<code>orchestrator</code>	Owns V2 active coordination state in Redis: dispatch, claim, leases, completion acceptance, recovery, and final state transitions.
<code>worker</code>	Performs Docker image builds and function execution. Maintains local warm containers and reports state.
<code>invocation-finalizer</code>	Commits staged invocation artifacts and moves accepted completions into visible terminal results.
<code>projector</code>	Projects orchestrator events into PostgreSQL so user-facing history remains queryable.
<code>cleaners</code>	Clean expired invocations, staged artifacts, dead-letter records, and registry images marked for deletion.

This split was not present at the very beginning. The first implementation was closer to a Django backend that directly queued and coordinated jobs. As the prototype moved toward multiple workers and reliability, the coordination logic was moved toward an orchestrator role.

3.2 Data stores

3.2.1 PostgreSQL

PostgreSQL stores persistent platform metadata. It is the long-term record for functions, versions, build attempts, invocation records, output metadata, log artifact pointers, ownership, and history projections. It is not the fastest possible coordination store, but it is reliable and easy to query.

3.2.2 Redis

Redis is used for active job coordination. The current V2 path uses Redis Streams and Redis keys for job state, worker leases, dispatch attempts, finalization markers, recovery sets, and worker metrics. Redis is faster than PostgreSQL for small coordination transitions and queue operations. It also allows the orchestrator to process jobs without performing a PostgreSQL query for every small state change.

3.2.3 Object storage

Media artifacts are stored through Django's storage interface. The deployed prototype now uses Parspack S3-compatible object storage. This is used for source bundles, invocation inputs, output files, logs, and invocation ZIPs. Before object storage was enabled on the deployed control plane, artifacts were stored in the local backend media volume; old artifacts were not migrated.

Object storage is not a finalization-only dependency. It appears in several parts of the lifecycle:

1. During function upload, the backend stores the source bundle.
2. During asynchronous invocation, the backend stores uploaded input files.
3. During execution, declared outputs may be uploaded directly by the runner or indirectly by the worker as staged artifacts.
4. During finalization, the backend verifies staged artifact metadata and marks the accepted files as committed.
5. During result download, the backend reads committed files and logs to produce the invocation ZIP.

The orchestrator does not read or write object bytes. It only stores and checks metadata: job IDs, attempts, completion IDs, terminal state, and artifact commit IDs. This is why object-storage failures should surface as artifact staging or finalization failures, not as scheduler-placement failures.

3.3 Why a separate userservice exists

The userservice was separated because authentication should not be tied to the load of function invocation and build traffic. If the serverless backend is busy, slow, or being restarted, users should still be able to log in, refresh tokens, and manage their account. The userservice signs RS256 JWTs. The backend verifies them locally by using the userservice public key or JWKS endpoint, so normal backend requests do not require a network call back to the userservice.

4 Job Model and Orchestration

4.1 Why orchestration became necessary

The earliest design could push work straight from the backend to Redis and let workers consume it. That works for a tiny prototype, but it becomes fragile as soon as the system needs multiple workers, retries, recovery, cancellation, fairness between builds and invocations, and visibility into queue state.

The orchestrator exists because the backend should behave more like a product API gateway and less like the active scheduler. The orchestrator owns the fast-changing coordination state. It chooses workers, grants leases, fences attempts, accepts completions, and recovers abandoned jobs.

4.2 Coordination versions

The project went through a V1 and V2 coordination model.

- | | |
|-----------|--|
| V1 | PostgreSQL job rows were the main coordination authority. Redis Lists carried job IDs to queues. Dispatch, claim, running, recovery, and final reports passed through the backend. |
| V2 | Redis Streams and orchestrator-owned state became the active authority. PostgreSQL remains a durable read model and history store through projection. |

V1 is deprecated. It is kept only as compatibility code during the transition. New development should focus on V2.

4.3 V2 job state

A V2 job follows a fenced state machine. The important states are:

- **queued**: job exists and is ready for placement.
- **dispatched**: orchestrator assigned it to a worker and attempt.
- **running**: worker successfully claimed the delivery.
- **recovering**: lease expired or delivery needs repair.
- **finalizing**: worker has finished and completion was accepted, but artifacts are not yet committed.

- `succeeded`, `failed`, `cancelled`, `dead_lettered`: terminal states.

The `dispatch_attempt` number is a fencing token. If a worker reports completion for an old attempt after the job was recovered and redispached, the orchestrator rejects the stale report.




Figure 4.1: Simplified V2 job state machine.

4.4 Why Redis Streams

Redis Streams were chosen instead of Pub/Sub because messages must survive temporary consumer failure. Pub/Sub is transient: if a consumer is offline when a message is published, the message is gone. Streams keep messages until they are acknowledged or trimmed. Consumer groups also expose pending entries, which lets the orchestrator and background services recover work after crashes.

4.5 Backoff and dead-letter behavior

Recovery is not allowed to retry forever. If a job keeps being redispached and recovered, it eventually becomes `dead_lettered`. This protects the platform from jobs that repeatedly fail because of a broken worker, bad image, or persistent infrastructure issue.

Backoff delays are different for builds and invocations. Builds are expensive and users tolerate a longer delay, so build recovery can wait longer between attempts. Invocations are more user-visible, so their retry delays start faster. The prototype uses this policy direction because it matches the product expectation: invocation latency matters more than build latency, while build storms can damage worker capacity.

5 Build Flow

5.1 Build overview

The build path converts user source into a runnable Docker image. A function can have only one active build in the current product semantics. If the user changes source code, the new source creates a replacement build. The old active image continues serving until the new build succeeds. If the new build fails, the existing active image stays active.




Figure 5.1: Build job flow.

5.2 Source bundle

The uploaded source bundle contains at least:

- `handler.py`
- `requirements.txt`
- `config.json`

The backend validates the bundle before storing it. Unsafe paths are rejected. That means entries such as absolute paths or parent-directory traversal are not allowed. This protects the builder from source bundles that try to write outside the expected build context.

5.3 Builder output

The worker builder generates a small runtime wrapper and Dockerfile. The runner imports the handler, loads the event, calls the function, writes the JSON result, and captures stdout/stderr through Docker logs and artifact files.

Listing 5.1: Simplified runner behavior.

```
1 handler_path = os.environ["FUNCTION_HANDLER"] # for example: handler.main
2 module_name, function_name = handler_path.rsplit(".", 1)
3
4 module = importlib.import_module(module_name)
5 handler = getattr(module, function_name)
6
7 with open(os.environ["FUNCTION_EVENT_PATH"], "r", encoding="utf-8") as f:
8     event = json.load(f)
9
10 context = {
11     "request_id": os.environ.get("FUNCTION_REQUEST_ID"),
12     "input_dir": os.environ.get("FUNCTION_INPUT_FILES_DIR"),
13     "output_dir": os.environ.get("FUNCTION_OUTPUT_DIR"),
14 }
15
16 result = handler(event, context)
17
18 with open(os.path.join(context["output_dir"], "result.json"), "w") as f:
19     json.dump(result, f)
```

The real runner contains more detail, including timing measurements, direct output upload support, and reserved file handling. The snippet above captures the core idea.

5.4 Build image lifecycle

Build output images are stored in a local Docker registry. In the current distributed deployment, the control plane hosts the registry and workers pull from it over the private network. Function image records track whether an image is active, superseded, failed, pending deletion, deleted, or delete-failed. This ledger matters because image deletion must not accidentally remove the active image for a function.

5.5 Build cancellation and retry

Build jobs support queued cancellation and cooperative running cancellation. The admin can configure a build retry policy. Per-attempt build history is stored so that a failed build can be examined without losing the overall build state. The current prototype has the policy mechanics; production-level build resource isolation and package scanning remain future work.

6 Invocation Flow in Detail

6.1 Asynchronous invocation, step by step

An asynchronous invocation is the most complete and important path in the system. It is designed for JSON events, optional input files, declared output files, durable polling, and artifact download.

1. The caller sends `POST /api/functions/{id}/invoke/`.
2. The backend authenticates the caller through JWT, function token, or public access.
3. The backend selects the active built function version.
4. The backend validates the JSON event and any input files.
5. Input files are stored as invocation input artifact records.
6. The backend creates an invocation record and a durable job creation event.
7. The outbox relay publishes the creation event to Redis.
8. The orchestrator reads the event and places the job on a worker stream.
9. The worker consumes the delivery and asks the orchestrator to claim it.
10. The orchestrator verifies the worker name, attempt, state, and lease.
11. The worker fetches payload and input files from protected backend endpoints.
12. The worker prepares a sandbox and runs or reuses a function container.
13. The runner loads the event and calls the handler.
14. The finalizer asks the backend to commit the staged artifacts.
15. The backend makes artifacts visible only after verification and commit.
16. The orchestrator records terminal state and emits a projection event.
17. The projector updates PostgreSQL history.
18. The caller polls the invocation endpoint and sees the final result.




Figure 6.1: Asynchronous invocation execution path. The backend owns artifact storage. The orchestrator owns Redis coordination. The worker runs the function and reports completion.

6.2 Invocation execution diagram

The important boundary is that artifact bytes do not pass through the orchestrator. Inputs, staged outputs, logs, and ZIP material are handled by backend-owned artifact APIs and object storage. The orchestrator only sees job metadata, attempts, leases, completion IDs, terminal status, and manifests.

6.3 Finalization and projection diagram




Figure 6.2: Finalization and projection path. The orchestrator does not write PostgreSQL directly; it emits Redis projection events consumed by the backend projector.

The finalizer is not the only component that can cause object-storage traffic. It is shown here only because it is responsible for committing already staged artifacts after the worker has finished executing. Earlier storage operations, such as source uploads and input uploads, are handled by the backend before the finalizer becomes involved.

6.4 Input sandbox

Input files are not simply mounted from arbitrary user paths. The backend stores uploaded input files, and the worker downloads them through internal endpoints. The worker writes

sanitized copies into a temporary sandbox. The function receives a directory path and a metadata JSON value so that it can open the files explicitly.

This design has overhead: files are uploaded to the backend, persisted, fetched by the worker, and copied into the sandbox. The tradeoff is that the platform owns validation, retention, download behavior, and access control.

6.5 Output sandbox

The function writes declared output files into `/sandbox/output`. The prototype has experimented with several strategies for collecting outputs:

- normal Docker volumes,
- tmpfs output mounts with size limits,
- Docker export/copy after execution,
- direct runner upload to the backend.

The current direction is performance-oriented: if no declared output files exist, output export is skipped. This matters because Docker archive and copy operations were visible in the timing reports even when the user function was small. Avoiding that work on the common JSON-only path is one of the simplest latency wins.

For output-producing invocations, the current implementation has two paths. The preferred fast path is direct runner upload. The worker starts the function container with scoped upload metadata in environment variables: upload URL, upload token, job ID, dispatch attempt, completion ID, declared output names, and byte limits. That data is not general backend access. It is permission to upload staged outputs for exactly one fenced execution attempt.

The following excerpt is from `worker/builder.py`. This is more useful than pseudocode because it shows the exact boundary between “a function wrote a file” and “the platform accepts this as a staged output”. The runner refuses to use direct upload unless every fencing field is present, reads the declared contract from the environment, scans the output directory, uploads accepted files, and returns a manifest. If the runner cannot produce a manifest for a function that declares outputs, the worker treats the invocation as failed.

Listing 6.1: Direct runner upload gate from `worker/builder.py`.

```
1 def upload_declared_outputs_directly(output_dir, timings, env=None):
2     if output_dir is None or not truthy(env_get(env, "
3         FUNCTION_OUTPUT_DIRECT_UPLOAD_ENABLED")):
4         return None
5
6     upload_url = env_get(env, "FUNCTION_OUTPUT_UPLOAD_URL", "")
7     upload_token = env_get(env, "FUNCTION_OUTPUT_UPLOAD_TOKEN", "")
8     job_id = env_get(env, "FUNCTION_OUTPUT_UPLOAD_JOB_ID", "")
9     completion_id = env_get(env, "FUNCTION_OUTPUT_UPLOAD_COMPLETION_ID", "")
10    dispatch_attempt = int_env("FUNCTION_OUTPUT_UPLOAD_DISPATCH_ATTEMPT", 0,
11                                env=env)
12    timeout = int_env("FUNCTION_OUTPUT_UPLOAD_TIMEOUT_SECONDS", 10, env=env)
```

```

11     if not upload_url or not upload_token or not job_id or not completion_id or
        dispatch_attempt < 1:
12         raise RuntimeError("Direct output upload is enabled but upload metadata
            is incomplete.")
13
14     scan_started = time.perf_counter_ns()
15     declared = json.loads(env_get(env, "FUNCTION_DECLARED_OUTPUT_FILES", "[]"
        or "[]"))
16     output_files = collect_declared_output_files(
17         output_dir,
18         declared,
19         int_env("FUNCTION_OUTPUT_MAX_FILES", len(declared), env=env),
20         int_env("FUNCTION_OUTPUT_MAX_FILE_SIZE_BYTES", 10 * 1024 * 1024, env=
            env),
21         int_env("FUNCTION_OUTPUT_MAX_TOTAL_SIZE_BYTES", 25 * 1024 * 1024, env=
            env),
22     )
23     timings["runner_output_scan_ms"] = elapsed_ms(scan_started)
24
25     upload_started = time.perf_counter_ns()
26     manifest = [
27         upload_file(upload_url, upload_token, job_id, dispatch_attempt,
28                     completion_id, item, position, timeout)
29         for position, item in enumerate(output_files)
30     ]
31     timings["runner_output_upload_ms"] = elapsed_ms(upload_started)
32     return manifest

```

The worker then decides whether Docker filesystem export is needed at all. This decision is one of the main performance details in the prototype. JSON-only functions avoid export completely. Direct-upload functions also avoid export because the runner already uploaded staged outputs from inside the container. Only the conservative fallback path copies `/sandbox/export` out of Docker and performs worker-side validation/upload.

Listing 6.2: Worker output path selection from worker/executor.py.

```

1 declares_outputs = self._declares_output_files(job)
2 if direct_output_upload:
3     result = self._read_result(None, stdout)
4     output_manifest = self._read_runner_output_manifest(stdout)
5     if output_manifest is None and (job.get("declared_output_files") or []):
6         return ExecutionResult(
7             status="failed",
8             error_message="Runner did not report direct output uploads.",
9             ...
10        )
11     output_manifest = output_manifest or []
12     timing["docker_export_copy_ms"] = 0
13     timing["output_validation_ms"] = 0
14     timing["output_upload_ms"] = 0
15 elif not declares_outputs:
16     result = self._read_result(None, stdout)
17     output_manifest = []

```

```

18     timing["docker_export_copy_ms"] = 0
19     timing["output_validation_ms"] = 0
20     timing["output_upload_ms"] = 0
21 else:
22     self._copy_directory_from_container(container, SANDBOX_EXPORT_PATH,
    export_dir)
23     output_files = self._validate_declared_output_files(job,
    effective_output_dir)
24     output_manifest = self._upload_output_files(job, output_files)

```

The backend still remains the authority for visibility. Direct runner upload does not mean “the user can now download the file”. The upload endpoint first verifies the runner upload token and the fenced attempt headers, then writes the file as **staged**. Staged files physically exist but are not returned by normal user-facing APIs.

Listing 6.3: Direct staged-output upload endpoint from backend/apps/invocations/views.py.

```

1 token = request.headers.get("X-Runner-Upload-Token", "")
2 job_id = request.headers.get("X-Job-Id", "")
3 dispatch_attempt = int(request.headers.get("X-Dispatch-Attempt", "0") or 0)
4 completion_id = request.headers.get("X-Completion-Id", "")
5 original_path = request.headers.get("X-Original-Path", "")
6
7 verify_runner_output_upload_token(
8     token,
9     request_id=request_id,
10    job_id=job_id,
11    dispatch_attempt=dispatch_attempt,
12    completion_id=completion_id,
13 )
14
15 uploaded_file = ContentFile(request.body or b"")
16 uploaded_file.name = original_path or "output"
17 uploaded_file.content_type = request.headers.get("Content-Type", "") or ""
18
19 output = store_staged_invocation_output_file(
20     invocation=invocation,
21     job_id=job_id,
22     dispatch_attempt=dispatch_attempt,
23     completion_id=completion_id,
24     uploaded_file=uploaded_file,
25     original_path=original_path,
26     checksum_sha256="",
27     position=position,
28 )

```

The final commit step is where staged artifacts become published artifacts. The backend verifies that the manifest reported by the worker/orchestrator matches the staged files already stored for that completion. Only then does it mark the completion and output files as **committed**. The serializer later filters on this committed state, which is why users cannot see half-finished outputs.

Listing 6.4: Backend staged-artifact commit boundary from backend/apps/invocations/services.py.

```

1 files = list(completion.output_files.order_by("position", "id"))
2 verify_output_manifest(files, output_manifest)
3
4 completion.terminal_status = terminal_status
5 completion.result = completion_payload.get("result") or {}
6 completion.stdout = log_preview(completion_payload.get("stdout", ""))
7 completion.stderr = log_preview(completion_payload.get("stderr", ""))
8 completion.exit_code = completion_payload.get("exit_code")
9 completion.output_manifest = output_manifest
10 completion.artifact_commit_id = completion.artifact_commit_id or uuid.uuid4()
11 completion.status = StagedCompletionStatus.COMMITTED
12 completion.committed_at = timezone.now()
13 completion.save()
14
15 completion.output_files.update(status=StagedCompletionStatus.COMMITTED)

```

6.6 Why staged artifacts exist

Without staging, a bad failure order could expose partial output files before the platform marks the invocation complete. Staging solves this. The worker uploads outputs as invisible staged artifacts tied to one invocation, one dispatch attempt, and one completion ID. Users cannot download them yet. Only after the orchestrator accepts completion and the finalizer commits artifacts do they become visible.




Figure 6.3: Staged artifact visibility rule.

6.7 Simple and advanced response modes

The default response mode is simple. Simple mode returns only the user-facing information: status, result, and output paths. Advanced mode returns the richer metadata already needed by debugging tools and the frontend.

Listing 6.5: Simple terminal invocation response.

```

1 {
2   "status": "succeeded",
3   "result": {"message": "Hello, Ilya!"},
4   "outputs": [
5     {
6       "path": "report.txt",
7       "download_url": "/api/invocations/9273/outputs/1480/download/"
8     }
9   ]

```

10 }

7 The Worker and Function Containers

7.1 Worker responsibilities

Workers are deliberately kept relatively simple. They do not own user accounts or product metadata. Their job is to register, heartbeat, receive assigned work, claim it, execute it, upload artifacts, and acknowledge the delivery only after responsibility has safely moved forward.

Each worker has bounded concurrency. The important settings are:

- `WORKER_MAX_CONCURRENCY`: total jobs in flight.
- `WORKER_MAX_INVOCATION_CONCURRENCY`: invocation jobs in flight.
- `WORKER_MAX_BUILD_CONCURRENCY`: build jobs in flight.

The worker prioritizes invocations over builds. This matters because builds can take much longer and should not cause short function calls to sit behind a large dependency installation.

7.2 Worker heartbeat

Workers send heartbeats containing worker identity, status, active job counts, queue names, and warm-container inventory. The scheduler/orchestrator uses this information to avoid dead workers, avoid workers with active builds for invocations, and route warm invocations when possible.

Listing 7.1: Simplified worker heartbeat data.

```
1 {  
2   "worker": "worker-1",  
3   "status": "online",  
4   "active_jobs": 3,  
5   "active_invocations": 3,  
6   "active_builds": 0,  
7   "max_invocation_concurrency": 12,  
8   "warm_containers": [  
9     {"function_version_id": 87, "idle": true, "last_used": 1788332000}  
10  ]  
11 }
```


7.3 Container create versus start

Docker container creation and container start are separate phases. Creation constructs a stopped container: filesystem layers are prepared, mounts are attached, limits are configured, environment variables are set, and metadata is registered with Docker. Starting the container begins execution of the process inside the already-created container.

This distinction matters in benchmarks. Cold invocations pay both create and start overhead. Warm containers avoid most repeated create/start cost, but they must be reset carefully enough to avoid leaking state between invocations.

7.4 Warm containers

Warm containers were added because cold container execution was too expensive for repeated small functions. The current warm-container design has these stages:

- reactive warm containers only,
- one warm container per function version per worker by default,
- TTL-based eviction,
- scheduler-aware exact warm routing,
- guarded sticky routing as a fallback,
- usage- and pressure-aware eviction.

The most important performance improvement came from moving to a resident warm runner. The old warm path still executed `python /runner.py` on each call. The resident path starts a small runner server inside the warm container and sends local requests to it. This avoids repeated Python startup and repeated module imports for no-op or small functions.

7.5 Warm container limitations

A warm container for function version A cannot safely be used for function version B. The image, dependencies, code, handler, environment, and possibly module-level state are all specific to one function version. Reusing a warm container across different functions would break isolation and correctness.

Warm containers also increase memory pressure. Keeping many containers alive improves latency for repeated calls but consumes worker resources. That is why the prototype uses TTL and pressure-aware eviction instead of keeping all warm containers forever.

TTL means *time to live*. In the warm-container pool, TTL is not a network TTL. It is the maximum amount of time an idle warm container is allowed to sit unused before the worker retires it. If a function is called repeatedly, the container's `last_used_at` timestamp is refreshed on release. If nobody uses it for `WORKER_WARM_IDLE_TTL_SECONDS`, the worker destroys it. In the local benchmark profile this was tuned lower for faster feedback; in a normal deployment it can be higher to keep frequently used functions warm.

The warm pool also has non-time-based retirement rules. A container can be retired because it is too old, because it has served too many invocations, or because a failed execution made it unsafe to reuse. These rules are deliberately separate from TTL:

- idle TTL answers: “has this idle container been unused too long?”
- maximum age answers: “has this container existed too long overall?”
- maximum uses answers: “has this container been reused enough times that refreshing it is safer?”
- failed execution answers: “did the last run leave this container in an unknown state?”

Pressure-aware eviction is the part that decides what to remove when the pool has too many containers. The pressure can come from three places: too many idle containers for one function version, too many containers in the whole pool, or too much warm-container memory reserved on that worker. This is better than plain FIFO because FIFO might delete a useful hot container while keeping a rarely used one. The implementation keeps per-container usage metadata and chooses victims using hit count, use count, recency, age, and memory size.

Listing 7.2: Warm-container retirement checks from worker/warm_pool.py.

```

1 def _retirement_reason_locked(self, record):
2     now = self._now()
3     if self.idle_ttl_seconds and now - record.last_used_at >= self.
idle_ttl_seconds:
4         return "idle_ttl"
5     if self.max_age_seconds and now - record.created_at >= self.max_age_seconds
:
6         return "max_age"
7     if self.max_uses and record.use_count >= self.max_uses:
8         return "max_uses"
9     return ""

```

Listing 7.3: Pressure-aware warm-pool victim selection from worker/warm_pool.py.

```

1 def _choose_eviction_victim_locked(self, candidates, *, pressure):
2     if pressure == "pool_memory_pressure":
3         return min(candidates, key=lambda record: (
4             record.hit_count,
5             record.use_count,
6             -record.key.memory_mb,
7             record.last_used_at,
8             record.created_at,
9         ))
10    return min(candidates, key=lambda record: (
11        record.hit_count,
12        record.use_count,
13        record.last_used_at,
14        record.created_at,
15        -record.key.memory_mb,
16    ))

```

The memory-pressure ordering intentionally treats large low-hit containers as good eviction candidates. Under normal container-count pressure, the pool first tries to remove low-hit, low-use, older idle containers. This is still a simple policy, but it captures the important intuition: keep containers that are actually being reused, and reclaim containers that are idle, cold in practice, or expensive relative to their benefit.

8 Scheduling Policy

8.1 What the orchestrator actually owns

The current service is named `scheduler` in the repository, but its V2 role is closer to an orchestrator. It is responsible for fast-changing active coordination state. It consumes V2 job creation events from Redis Streams, creates Redis job-state hashes, places ready jobs onto worker-specific streams, validates worker claims, grants leases, accepts completions, starts finalization for invocation artifacts, marks build jobs terminal, recovers expired leases, and emits projection events for PostgreSQL read models.

The orchestrator does not own source files, input files, output files, user accounts, or function metadata. Those stay behind the backend and userservice. The orchestrator owns the execution state machine. That distinction is important: the backend is the product API; Redis plus the orchestrator are the active control plane for V2 jobs.

8.2 Job creation and ready set

The backend creates a durable `job.created` event. The orchestrator consumer reads that event from Redis, ignores non-V2 events, creates a Redis job state, stores the payload, and inserts the job into a sorted ready set keyed by `available_at_ms`. Delayed retries use the same ready set. A job is only eligible for dispatch when its score is less than or equal to the current time.

Listing 8.1: V2 job creation ingestion from `scheduler/production_orchestrator.py`.

```
1 def process_creation_event(self, stream_id, fields):
2     if fields.get("event_type") != "job.created":
3         return False
4     event = json.loads(fields["payload"])
5     if int(event.get("coordination_version", 1)) != 2:
6         return False
7
8     job_id = str(event["job_id"])
9     available_at_ms = timestamp_ms(event.get("available_at"))
10    self.state.create_job(job_id, event["job_type"], available_at_ms=
available_at_ms)
11    self.redis.hset(self.state.key(job_id), mapping={
12        "payload": json.dumps(event.get("payload") or {}, separators=(",", ":"))
13    },
14        sort_keys=True),
15        "source_event_id": fields.get("event_id", ""),
16        "source_stream_id": stream_id,
```

```
16         "max_recovery_attempts": int(event.get("max_recovery_attempts", 3)),
17     })
18     self.redis.zadd(self.ready_key, {job_id: available_at_ms})
19     return True
```

8.3 Worker heartbeats as scheduling input

Workers periodically report their current state. The orchestrator stores each heartbeat in Redis and puts the worker name into a lease sorted set. If a worker misses its heartbeat lease, `list_online_workers` expires it before placement. This means the placement policy only sees workers that are currently online according to Redis time and the configured worker lease TTL.

Listing 8.2: Worker heartbeat fields from scheduler/orchestrator_workers.py.

```
1 redis.call('HSET', KEYS[1],
2     'name', ARGV[1],
3     'hostname', ARGV[2],
4     'status', ARGV[16],
5     'max_concurrency', ARGV[3],
6     'max_build_concurrency', ARGV[4],
7     'max_invocation_concurrency', ARGV[5],
8     'active_jobs', ARGV[6],
9     'active_builds', ARGV[7],
10    'active_invocations', ARGV[8],
11    'metadata', ARGV[13],
12    'last_seen_ms', ARGV[14],
13    'lease_expires_at_ms', ARGV[15])
```

The `metadata` field includes warm-pool inventory. That is how the orchestrator learns that a worker has, for example, an idle resident warm container for function version 87 with a specific image, handler, memory limit, and output tmpfs size. The orchestrator does not inspect Docker directly. It trusts worker heartbeats as its worker-capacity view.

8.4 Separate build and invocation streams

V2 separates build and invocation deliveries. Each worker has a build stream and an invocation stream. This prevents long build jobs from sitting directly in front of latency-sensitive invocations. The worker loop also prioritizes invocation stream reads when invocation capacity exists.

- job creation stream: backend-to-orchestrator events,
- ready sorted set: orchestrator-owned dispatch eligibility,
- worker invocation stream: V2 invocation deliveries for one worker,
- worker build stream: V2 build deliveries for one worker,
- finalization stream: invocation artifact commit work,
- projection stream: Redis terminal/running events copied to PostgreSQL.

8.5 Invocation placement policy

Invocation placement is more sophisticated than round robin. It uses reported worker state plus short-lived local reservations inside the orchestrator process. The local reservations matter because heartbeats are periodic. If the orchestrator dispatches five jobs to the same worker between two heartbeats, the worker’s last reported `active_jobs` may still look low. Local reservations prevent the orchestrator from overbooking a worker before the next heartbeat arrives.

The current invocation order is:

1. remove expired local dispatch reservations, warm reservations, and sticky routes;
2. compute the exact warm-container key for the invocation;
3. prefer online workers with no active or queued builds, available invocation capacity, and an exact idle warm container;
4. if several warm workers match, choose the lowest invocation load and round-robin among ties;
5. if no warm worker is visible, consider workers without active or queued builds;
6. filter to workers below `max_invocation_concurrency` when possible;
7. try guarded sticky routing only if the recently used worker is still capacity-safe;
8. otherwise choose the least-loaded candidate and round-robin among ties.

Listing 8.3: Warm-first invocation placement from scheduler/scheduler.py.

```

1 warm_candidates = [
2     worker
3     for worker in workers
4     if worker_active_builds(worker) == 0
5     and worker_queued_builds(worker) == 0
6     and invocation_load(worker, local_invocation_loads=local_invocation_loads,
7                          local_load_ttl_seconds=local_load_ttl_seconds,
8                          now=now) < worker_max_invocation_concurrency(worker)
9     and worker_idle_warm_count(worker, warm_key,
10                                local_warm_reservations=local_warm_reservations,
11                                local_load_ttl_seconds=local_load_ttl_seconds,
12                                now=now) > 0
13 ]
14 if warm_candidates:
15     best_load = min(invocation_load(worker, local_invocation_loads=
16     local_invocation_loads,
17                                local_load_ttl_seconds=
18                                local_load_ttl_seconds,
19                                now=now)
20                        for worker in warm_candidates)
21     best = [worker for worker in warm_candidates
22             if invocation_load(worker, local_invocation_loads=
23             local_invocation_loads,

```

```

21         local_load_ttl_seconds=local_load_ttl_seconds,
22         now=now) == best_load]
23     return choose_round_robin_worker(best, "warm-invocation", round_robin_state
    )

```

The warm key is intentionally exact. A warm container is only considered a match when function version, image reference, handler, memory size, and output tmpfs size match. A container for one function version cannot safely serve another function version.

Listing 8.4: Exact warm-container match from scheduler/scheduler.py.

```

1 return (
2     str(item.get("function_version_id") or "")
3     == str(warm_key.get("function_version_id") or "")
4     and str(item.get("image_ref") or "") == str(warm_key.get("image_ref") or "")
5     and str(item.get("handler") or "") == str(warm_key.get("handler") or "")
6     and safe_int(item.get("memory_mb")) == safe_int(warm_key.get("memory_mb"))
7     and safe_int(item.get("output_tmpfs_size_bytes"))
8     == safe_int(warm_key.get("output_tmpfs_size_bytes"))
9 )

```

8.6 Guarded sticky fallback

Guarded sticky routing is a prediction layer, not the primary policy. It exists because a worker may have created a warm container after a recent invocation, but the next heartbeat may not have reported that warm container yet. The orchestrator therefore remembers the recent route for a function version for a short TTL, currently ten seconds in the policy path. It only follows that route when the worker still has no build pressure, still has invocation capacity, and is not much more loaded than the best reported candidate.

Listing 8.5: Guarded sticky checks from scheduler/scheduler.py.

```

1 best_reported_load = min(worker_reported_invocation_load(worker)
2                           for worker in candidates)
3 for worker in candidates:
4     if worker.get("name") != worker_name:
5         continue
6     if worker_active_builds(worker) != 0 or worker_queued_builds(worker) != 0:
7         return None
8     reported_load = worker_reported_invocation_load(worker)
9     local_dispatches = local_invocation_load(worker,
10        local_invocation_loads=local_invocation_loads,
11        local_load_ttl_seconds=local_load_ttl_seconds,
12        now=now)
13     if local_dispatches > predictive_sticky_max_local_dispatches:
14         return None
15     if reported_load + local_dispatches >= worker_max_invocation_concurrency(
16        worker):
17         return None
18     if reported_load > best_reported_load + predictive_sticky_load_slack:
19         return None

```

```
19 return worker
```

The important parameters are not arbitrary adjectives. “Capacity-safe” means reported active work plus locally remembered dispatches is below `max_invocation_concurrency`. “Not much more loaded” means the sticky worker’s reported load is no more than `predictive_sticky_load_slack` above the best candidate’s reported load. “Local burst limit” means the orchestrator will not send more than `predictive_sticky_max_local_dispatches` recent unconfirmed local dispatches to the same sticky worker.

8.7 Build placement policy

Build placement is intentionally stricter than invocation placement. The build selector only chooses workers with zero active jobs, zero queued invocations, and zero queued builds. This protects the user-visible invocation path from dependency installation and Docker build pressure. It can reduce build throughput, but it gives the prototype a clearer latency priority.

Listing 8.6: Build placement from `scheduler/scheduler.py`.

```
1 def choose_build_worker(workers, *, round_robin_state):
2     candidates = [
3         worker
4         for worker in workers
5         if worker_active_jobs(worker) == 0
6         and worker_queued_invocations(worker) == 0
7         and worker_queued_builds(worker) == 0
8     ]
9     if not candidates:
10         return None
11     return choose_round_robin_worker(candidates, "build", round_robin_state)
```

8.8 Atomic dispatch

After policy chooses a worker, dispatch must be atomic. It is not enough to set `dispatched` and later push to a worker stream in separate operations, because the orchestrator could crash between those operations. The production orchestrator uses a Redis Lua script that checks the job is still queued, increments `dispatch_attempt`, writes assigned worker and lease metadata, adds the delivery to the worker stream, stores the stream ID, and removes the job from the ready set in one Redis operation.

Listing 8.7: Atomic dispatch script from `scheduler/production_orchestrator.py`.

```
1 local attempt = redis.call('HINCRBY', KEYS[1], 'dispatch_attempt', 1)
2 redis.call('HSET', KEYS[1],
3     'status', 'dispatched',
4     'assigned_worker', ARGV[1],
5     'worker_stream', KEYS[3],
6     'effective_image_ref', image_ref,
7     'lease_expires_at_ms', ARGV[6],
8     'updated_at_ms', ARGV[4])
```



```

9 local stream_id = redis.call('XADD', KEYS[3], '*',
10     'job_id', ARGV[2],
11     'dispatch_attempt', tostring(attempt),
12     'job_type', ARGV[3])
13 redis.call('HSET', KEYS[1], 'delivery_stream_id', stream_id)
14 redis.call('ZREM', KEYS[2], ARGV[2])
15 return {'ok', 'dispatched', tostring(attempt), stream_id}

```

For build jobs, the dispatch script also creates an attempt-specific image tag. The active function version is not switched to that image until the build completes successfully and the terminal projection is committed. This prevents a failed replacement build from breaking the currently working image.

8.9 Claim and lease protocol

Dispatch does not mean the worker may execute blindly. The worker reads a delivery from its Redis Stream and calls the orchestrator claim API with `worker_name` and `dispatch_attempt`. The orchestrator atomically checks that the job is still assigned to that worker, that the attempt matches, and that the job is still in `dispatched`. If the check passes, the job moves to `running`, a lease expiry is written, and the orchestrator returns the execution payload. If the check fails, the worker acknowledges the stale delivery and does not execute it.

Listing 8.8: Worker-side V2 claim before execution from `worker/worker.py`.

```

1 claim = orchestrator.claim_job(job_id, {
2     "worker_name": worker_name,
3     "dispatch_attempt": int(delivery.get("dispatch_attempt", 0)),
4 })
5 if not claim.get("claimed"):
6     client.xack(worker_stream, V2_WORKER_GROUP, stream_id)
7     continue
8
9 job = dict(claim.get("payload") or {})
10 pool.submit(run_v2_claimed_job, job=job, backend=backend,
11             orchestrator=orchestrator, worker_name=worker_name, ...)

```

Listing 8.9: Orchestrator claim transition from `scheduler/orchestrator_state.py`.

```

1 local worker = redis.call('HGET', KEYS[1], 'assigned_worker') or ''
2 local attempt = tonumber(redis.call('HGET', KEYS[1], 'dispatch_attempt') or '0'
3 )
4 if worker ~= ARGV[1] or attempt ~= tonumber(ARGV[2]) then
5     return {'stale', status, tostring(attempt)}
6 end
7 if status ~= 'dispatched' then
8     return {'invalid_state', status, tostring(attempt)}
9 end
10 redis.call('HSET', KEYS[1],
11     'status', 'running',
12     'lease_expires_at_ms', ARGV[3],
13     'updated_at_ms', ARGV[4])

```

```
13 return {'ok', 'running', tostring(attempt)}
```

The lease is renewed while the worker is executing. If the worker process dies, the worker machine loses network connectivity, or the job stops renewing its lease, the orchestrator can recover the job after the lease expires. The dispatch attempt fencing token prevents an old worker from successfully reporting completion after a newer attempt has already been created.

8.10 Completion, finalization, and ACK

For invocations, completion is split into two phases. First, the worker uploads staged outputs and reports completion to the orchestrator. The orchestrator verifies the worker and attempt, stores the completion payload in Redis, moves the job to **finalizing**, removes the running lease, ACKs the worker delivery, and emits a finalization event. At that point the function will not be executed again for that attempt; responsibility has moved to the finalizer.

Listing 8.10: Invocation completion acceptance from scheduler/production_orchestrator.py.

```
1 begun = self.state.begin_finalization(
2     job_id,
3     worker_name=worker_name,
4     dispatch_attempt=dispatch_attempt,
5     completion_id=completion_id,
6     completion_payload=completion_payload,
7     terminal_status=status,
8     now_ms=now_ms,
9 )
10 if job_before.get("job_type") == "invocation":
11     self.redis.zrem(self.lease_key, job_id)
12     self.ack_delivery(self.state.get_job(job_id))
13     self.emit_finalization(job_id)
14     return {"accepted": True, "status": V2JobStatus.FINALIZING, ...}
```

The finalizer later asks the backend to commit staged artifacts. After the backend returns an **artifact_commit_id**, the finalizer calls the orchestrator again. The orchestrator then moves **finalizing** to the terminal state and emits a projection event. This second transition checks that a successful invocation has an artifact commit ID; otherwise a succeeded result would become visible without committed artifacts.

Listing 8.11: Finalization guard from scheduler/orchestrator_state.py.

```
1 if status ~= 'finalizing' then
2     return {'invalid_state', status, tostring(attempt)}
3 end
4 if completion ~= ARGV[1] then
5     return {'completion_conflict', status, tostring(attempt)}
6 end
7 if ARGV[2] == 'succeeded' and ARGV[3] == '' then
8     return {'missing_artifact_commit', status, tostring(attempt)}
9 end
10 redis.call('HSET', KEYS[1],
```

```
11     'status', ARGV[2],
12     'artifact_commit_id', ARGV[3],
13     'finished_at_ms', ARGV[4],
14     'updated_at_ms', ARGV[4])
```

8.11 Recovery and dead letters

Recovery is driven by the orchestrator's job lease sorted set. When a dispatched or running job's lease expires, the orchestrator checks its state. If the job has not exceeded its maximum recovery count, it moves the job back to `queued`, sets a future `available_at_ms` according to the backoff policy, clears worker assignment fields, and redispaches it later. If it has exceeded the configured recovery count, it becomes `dead_lettered`. This keeps the platform from retrying a permanently broken job forever.

Listing 8.12: Recovery transition from scheduler/orchestrator_state.py.

```
1 if attempt ~= tonumber(ARGV[1]) then
2     return {'stale', status, tostring(attempt)}
3 end
4 local recovery = redis.call('HINCRBY', KEYS[1], 'recovery_count', 1)
5 redis.call('HSET', KEYS[1],
6     'status', 'queued',
7     'available_at_ms', ARGV[2],
8     'last_error', ARGV[3],
9     'updated_at_ms', ARGV[4])
10 redis.call('HDEL', KEYS[1],
11     'assigned_worker', 'worker_queue', 'lease_expires_at_ms',
12     'completion_id', 'completion_payload', 'completion_status',
13     'artifact_commit_id')
```

8.12 What can go wrong

Scheduling is a balancing act. Too much stickiness can overload one worker while others are idle. Too little stickiness causes repeated cold starts. Too much build parallelism slows invocations. Too little build parallelism makes builds wait unnecessarily. The current policy is practical but not final. Future work can add richer signals such as recent CPU load, Docker daemon pressure, image cache hits, per-function hotness, queue age, and per-worker p95 latency.

9 Security and Access Control

9.1 Authentication

Account authentication is handled by the userservice. It issues RS256 JWTs. The backend validates those tokens locally using the public key/JWKS exposed by the userservice. This avoids a network call for every authenticated backend request.

9.2 Authorization

The backend enforces ownership rules for functions, versions, builds, invocations, outputs, and token management. Admin users can access worker-node management APIs. Internal worker and service endpoints use a shared internal token through `X-Internal-Token`.

9.3 Invocation tokens

Invocation tokens are function-scoped opaque secrets. Creating or rotating a token returns the raw token once. Later API responses only show metadata such as token prefix, name, expiry, revoked state, and last-used time.

When an invocation is created by function token or public access, the platform returns an invocation read token. This read token lets that caller poll and download only the result of that specific invocation. A different token caller should not be able to read another token caller's invocation output.

9.4 Current hardening gaps

The prototype is not yet hardened as a hostile multi-tenant platform. Important future work includes CPU controls, process-count limits, read-only root filesystems, seccomp/AppArmor profiles, stricter network controls, dependency scanning, image scanning, and build isolation from the host Docker daemon.

10 Artifacts, Logs, and Retention

10.1 Artifact categories

The platform handles several artifact types:

- source bundles,
- invocation input files,
- invocation output files,
- stdout and stderr logs,
- invocation ZIP downloads,
- Docker images.

Source bundles are kept permanently according to the current project policy. Invocation inputs, outputs, logs, and invocation records expire after seven days. If an invocation expires, the API returns 404 as if the invocation never existed.

10.2 Log storage

Logs are not stored as large database text fields. The current design stores stdout and stderr as file artifacts, with database pointers and previews. This avoids making PostgreSQL carry potentially large log bodies while still giving the frontend enough information for status and basic display.

Full logs are included in the invocation ZIP. The user does not get separate log browsing endpoints in the prototype UI contract.

10.3 Object storage

The deployed platform uses Parspack S3-compatible object storage for media artifacts. The backend accesses it through Django storage, so most application code uses the same file API whether storage is local or S3-backed.

This was recently verified with a public-domain smoke test. A function was created, built, invoked, produced `report.txt`, and the backend confirmed that both the output file and stdout log used `S3Storage` and existed in the configured bucket.

11 Reliability Design

11.1 Fencing tokens

The platform uses dispatch attempts as fencing tokens. If a worker runs an old delivery after recovery, its completion report contains the wrong attempt and is rejected. This protects against stale workers overwriting newer state.

11.2 Leases and recovery

Dispatched and running V2 jobs hold leases. If the lease expires, recovery can requeue the job or move it toward dead-letter depending on the recovery count. The goal is to avoid permanent stuck jobs while also avoiding infinite retry loops.

11.3 Safe completion protocol

The safe completion protocol avoids cross-database transactions between Redis and PostgreSQL. Instead, it uses staging, idempotency, fencing, and retryable finalization.

1. Worker uploads staged outputs.
2. Worker reports completion with job ID, attempt, completion ID, status, result, and manifest.
3. Orchestrator verifies running state and attempt, stores completion payload, moves to **finalizing**, and acknowledges worker delivery.
4. Finalizer asks backend to commit staged artifacts.
5. Backend verifies and marks artifacts committed in a transaction.
6. Orchestrator verifies artifact commit ID and moves the job to terminal state.
7. Projector records terminal history in PostgreSQL.

The key rule is that a physical file may exist before the invocation is terminal, but it is not visible to users until the final terminal transition.

11.4 Graceful shutdown

Workers support draining behavior. A draining worker should stop accepting new work while letting in-flight jobs finish. This is necessary for maintenance, deployments, and removing workers from the distributed deployment without turning normal in-flight invocations into artificial failures.

12 Observability

12.1 Stack

The observability stack contains:

Prometheus Scrapes and stores numeric metrics.

Grafana Displays dashboards backed by Prometheus and Loki.

Loki Stores searchable logs.

Grafana Alloy
 Collects Docker logs from the control-plane host and ships them to Loki.

cAdvisor Exposes container CPU, memory, filesystem, and network metrics.

Redis exporter
 Exposes Redis metrics.

Postgres exporters
 Expose backend and userservice database metrics.

The Grafana dashboard shows scrape health, worker reachability, queue backlog, worker current state, queue pressure, worker execution, orchestrator stream lag, finalization health, duplicate/fencing signals, storage-related errors, and recent logs.

12.2 Current observability limitation

Loki currently collects control-plane Docker logs only. Worker metrics are scraped from the worker VMs, but worker container logs are not shipped yet. The next observability improvement should deploy Alloy or another log shipper on each worker machine so that worker timing and failure logs appear in Loki.

12.3 Interpreting empty panels

An empty error panel does not necessarily mean observability is broken. It can mean no error log line matched the filter in the selected time range. For this reason, the dashboard includes a general recent platform logs panel. If normal logs appear there but the error panel is empty, the system is observable and no matching recent errors were found.

13 Deployment

13.1 Local deployment

The local deployment remains available through `docker-compose.yml`. It is used for development and many of the repeatable tests.

Listing 13.1: Local stack startup.

```
1 docker compose up -d
2 docker compose exec backend python manage.py migrate
3 docker compose exec userservice python manage.py migrate
```

13.2 Distributed prototype deployment

The first distributed deployment uses one control-plane machine and two worker machines. The control plane hosts frontend, Nginx, userservice, backend, PostgreSQL, Redis, registry, orchestrator services, finalizer, projector, cleaners, and observability. Each worker machine runs a worker container and a local Docker Engine.

The current public domain is `fleeingcircus.ir`. Nginx terminates HTTPS with a Let's Encrypt certificate. Grafana is exposed through `/grafana/` and protected by Nginx basic authentication.

13.3 Private network

The deployment uses a VPC private network. Workers connect to the control plane over private IPs for Redis, backend internal APIs, orchestrator APIs, and the Docker registry. PostgreSQL is not exposed publicly. Redis and registry should also remain restricted to private worker IPs.

13.4 Registry behavior

The first distributed version uses a single registry on the control-plane machine. Workers pull images from that registry. This creates a pull delay when a worker does not have an image cached locally. Multiple registries were considered, but for the prototype a single reachable registry is simpler and good enough. Future systems can consider per-region registries or scheduler placement based on image cache locality.

14 Benchmarks and Performance

14.1 Benchmark terminology

Earlier benchmark reports used the word concurrency in a way that needed clarification. Two different ideas matter:

Saturation concurrency

The number of client-side invocation requests kept in flight by the benchmark script.

Worker invocation concurrency

The number of simultaneous invocation jobs each worker process may run.

The later reports renamed and clarified this distinction. Saturation concurrency tests show how the system behaves under client load. Worker concurrency tests isolate how much parallel work each worker should run.

14.2 Warm-container benchmark

Warm containers improved repeated invocations significantly. In the Stage 1 API-level sequential benchmark, the cold worker ran six no-op invocations with a client wall time of about 17.6 seconds and median invocation duration of about 1.62 seconds. The warm-enabled worker reduced client wall time to about 12.0 seconds and median invocation duration to about 0.925 seconds.

The bigger improvement came later from resident warm containers:

Pattern	Old wall	New wall	Old median	New median	Speedup
No-op	12.5 s	3.8 s	3918 ms	88 ms	44.52x
Sleep 3s	22.1 s	16.5 s	7266 ms	4997 ms	1.45x
No-op with output	11.9 s	9.1 s	3774 ms	1457 ms	2.59x

The interpretation is important. Warm containers help most when platform overhead dominates user code. If the user function sleeps for three seconds, there is less non-user work to remove proportionally.

14.3 Distributed worker concurrency benchmark

The most recent distributed benchmark used one control-plane VM and two worker VMs. Each test began from a clean platform state: no pending scheduler queues, no worker backlog, and no active invocations. The mixed workload contained four function styles:

- **tiny**: JSON-only function that returns immediately.
- **sleep**: JSON-only function that sleeps for one second.
- **dependency**: function importing and using a dependency built into the image.
- **input/output**: asynchronous function with an input file and a declared output file.

The table below isolates worker invocation concurrency. Saturation concurrency was held at 32 client-side in-flight requests. With two workers, cluster capacity is twice the per-worker value.

Scenario	Worker conc.	Cluster cap.	Avg wall	Avg thr.	P95 term.
mixed-30	2	4	13.12 s	2.31/s	12.34 s
mixed-30	4	8	11.24 s	2.67/s	10.09 s
mixed-30	8	16	10.34 s	2.91/s	9.63 s
mixed-30	16	32	11.47 s	2.64/s	10.59 s
mixed-80	2	4	32.13 s	2.50/s	12.95 s
mixed-80	4	8	24.34 s	3.29/s	11.17 s
mixed-80	8	16	24.23 s	3.30/s	12.16 s
mixed-80	16	32	24.54 s	3.26/s	14.02 s
mixed-320	2	4	117.27 s	2.73/s	15.36 s
mixed-320	4	8	96.44 s	3.32/s	14.41 s
mixed-320	8	16	91.44 s	3.50/s	13.17 s
mixed-320	12	24	93.33 s	3.43/s	13.83 s

The result is not linear scaling. Worker concurrency 2 underused the cluster. Moving to 4 improved throughput clearly. Moving to 8 helped the largest mixed-320 workload, but the smaller workloads were already close to saturation. Moving to 12 or 16 did not produce a clean win. At that point the workers were asking the local Docker daemons, disks, Python runtimes, artifact paths, and cleanup code to do too much at once. More in-flight work increased pressure without reducing wall time.

The two workers handled work reasonably evenly. In the mixed-320 concurrency 12 run, one repeat assigned 165 invocations to worker 1 and 155 to worker 2; the second assigned 172 and 148. This is not perfectly equal, but it shows that both workers were being used and that the orchestrator was not stuck on one machine.

Worker phase, mixed-320 c12	Median	P90	P95	Max
Docker image pull/check	333 ms	785 ms	993 ms	1742 ms
Sandbox preparation	175 ms	465 ms	510 ms	979 ms

Warm container create	1301 ms	2179 ms	2316 ms	2978 ms
Runner module imports	1004 ms	1601 ms	1750 ms	2561 ms
Warm runner execution	1618 ms	2721 ms	3253 ms	4441 ms
Docker export/copy	0 ms	0 ms	733 ms	1650 ms
Output upload	0 ms	0 ms	251 ms	999 ms
Warm sandbox cleanup before	355 ms	664 ms	766 ms	944 ms
Warm sandbox cleanup after	346 ms	608 ms	722 ms	1120 ms
Warm pool release	754 ms	1378 ms	1576 ms	2643 ms
Orchestrator claim	10 ms	16 ms	20 ms	34 ms
Orchestrator completion	11 ms	19 ms	22 ms	32 ms
Worker process total	5568 ms	7794 ms	8180 ms	10246 ms

This table is the strongest evidence about current bottlenecks. Orchestrator API calls are not the main cost; claim and completion are tens of milliseconds. The expensive phases are worker-local: container creation, Python imports, resident runner execution under load, Docker image checks/pulls, sandbox cleanup, and warm-pool release. Output work is mostly zero because many mixed workload calls do not declare output files, but the p95 values still show that artifact-producing calls pay a visible extra cost.

14.4 Invocation latency profiling before warm-container tuning

Before the resident warm runner existed, the project profiled the invocation path with worker-side timestamps. This was important because user-visible latency alone did not explain whether time was spent queueing, creating containers, importing Python modules, moving files, uploading artifacts, or reporting state.

Scenario	Wall	Queue	Worker	Overhead
No-op sequential	15.03 s	376 ms	1863 ms	1863 ms
One-second sequential	23.81 s	390 ms	2531 ms	1531 ms
Three-second concurrent burst	23.47 s	8589 ms	6502 ms	3502 ms

The concurrent burst showed two problems at once. First, queue-to-start time became high because the workers could not begin every invocation immediately. Second, work that did start also became slower because concurrent Docker operations contended with each other. A three-second user function had a median worker executor duration of about 6.5 seconds, meaning roughly 3.5 seconds were platform overhead.

Worker phase in concurrent three-second burst	Median
Sandbox preparation	419 ms
Docker volume create	75 ms
Docker container create	581 ms
Docker input copy	169 ms
Docker container start	564 ms
Docker wait, including user sleep	4390 ms

Docker logs read	138 ms
Docker export/copy	137 ms
Docker cleanup	202 ms
Backend running report	445 ms
Backend final report	172 ms
Executor duration	6502 ms
Worker process total	7300 ms

This profiling directly motivated three later changes: moving active V2 coordination out of backend report endpoints, adding warm containers, and skipping output export when no declared outputs exist.

14.5 Where time is currently spent

The performance reports show several recurring overhead sources:

- Docker image pulls when the worker does not have the image cached.
- Docker container create/start on cold paths.
- Python runtime startup and module import on non-resident paths.
- Docker copy/export when output files exist.
- Sandbox preparation and cleanup.
- Backend artifact upload and finalization for output-producing invocations.
- Polling delay from client-side or test-script status polling.

The resident warm runner is the biggest successful optimization so far for small JSON-only functions. For output-producing functions, artifact handling remains a visible bottleneck.

15 Things Tried, Deprecated, and Changed

15.1 Backend-owned coordination

The first path made the backend responsible for much of dispatch, claim, and status mutation. It was easier to build but less aligned with a distributed system. It also put fast-changing coordination logic on PostgreSQL. The current direction moves active coordination to the orchestrator and Redis while keeping PostgreSQL as persistent metadata and history.

The original shape was useful because it gave the project a working API quickly. A request could create a database row, the worker could call backend endpoints, and the backend could mutate the row from queued to running to terminal. The weakness appeared once we started thinking in distributed terms: the backend was becoming both the public product API and the scheduler/control loop for the execution system.

Old behavior	Problem	Replacement
Backend accepted worker claims	Public API service became a job coordinator	Orchestrator claim API with Redis state
Worker reports mutated PostgreSQL directly through backend endpoints	Fast-changing state created database pressure	Redis terminal state plus projector
Status was mostly PostgreSQL-backed	Active reads could lag behind real orchestration state	Redis active state, PostgreSQL history

The fix was not to remove PostgreSQL. PostgreSQL is still valuable for durable product metadata and history. The fix was to stop using it as the hot path for every small coordination transition. Redis holds active state; the projector writes a durable read model after the orchestrator emits projection events.

15.2 Redis Lists

Redis Lists worked for a simple queue but were less expressive for reliable consumer recovery and pending-entry inspection. Redis Streams are now preferred for V2 coordination because consumer groups and pending entries give better recovery mechanics.

Lists were good enough for the first worker stub: push a job into a list, pop a job from a list, execute it. The problem is what happens after a crash. If a worker destructively pops a job and dies before reporting completion, the queue itself no longer knows that

the job existed. We added processing queues and recovery logic, but this was building stream-like behavior manually.

Listing 15.1: Why Redis Streams replaced the list-only path.

```
1 Pub/Sub: fast broadcast, but messages disappear if nobody is listening.  
2 List: durable queue, but recovery metadata must be built manually.  
3 Stream: durable log with consumer groups and pending-entry inspection.
```

Streams fit the V2 protocol better because a delivery can be claimed, inspected, recovered, and acknowledged with clearer semantics. They also make it easier to run multiple workers without inventing a custom pending-delivery database.

15.3 Docker export/copy for outputs

The project tried Docker export/copy to avoid keeping tmpfs containers alive after function exit. This was correct but expensive under concurrent workload. The direct runner upload path was added to reduce Docker copy contention. Output export is skipped entirely when a function has no declared output files.

The tmpfs idea was attractive because the function would write into a bounded memory-backed filesystem. If the user wrote too much, the write would fail inside the function instead of filling the host disk. The catch was lifecycle: once the function container exited, the tmpfs contents were tied to the container. Keeping the container alive just so the worker could copy results out was wasteful. Exporting/copying after execution solved correctness but introduced Docker archive overhead.

The current compromise is:

1. If there are no declared outputs, skip output export entirely.
2. If direct runner upload is enabled, validate declared files inside the container and upload them directly to backend staging.
3. If direct upload is disabled or fails before completion reporting, fall back to worker-side validation and upload.

This does not make output-producing functions free. It only removes an avoidable Docker extraction step from the fast path.

15.4 Keeping containers alive with sleep

One idea was to keep a container alive after user code completed so that the worker could extract tmpfs outputs. This was rejected because it would hold CPU and memory resources after function completion and would not scale well with many requests. The project moved toward export/copy and then direct upload instead.

The rejection was specifically about scale. A 30-second sleep sounds harmless when one invocation is running. With 100 requests, it becomes 100 containers that have already finished useful work but still occupy runtime slots, memory, container metadata, and Docker daemon attention. Warm containers are different: they are intentionally pooled, bounded, reused, and evicted. A sleeping finished request is not a warm pool; it is delayed cleanup.

15.5 Sticky routing alone

Sticky routing was initially attractive because repeated calls to the same function are likely to benefit from the same worker. Blind stickiness can overload one worker. The current policy uses exact warm matches first and only uses guarded sticky routing as a fallback when it still looks capacity-safe.

The current scheduling preference is layered:

Listing 15.2: Current invocation placement policy, simplified.

```
1 if worker_has_idle_warm_container(function_version):  
2     choose_that_worker()  
3 elif recent_route_exists(function_version) and worker_still_has_capacity():  
4     choose_recent_worker()  
5 else:  
6     choose_least_loaded_worker_without_active_builds()
```

This keeps the good part of stickiness, namely likely cache and warm-container reuse, without letting old routing hints dominate the system after a worker becomes busy.

15.6 Single-service authentication

The backend originally contained account authentication as an app. The project then moved to a separate userservice because account login should not be coupled to serverless execution load. The backend still has compatibility bridge code, but the intended direction is external userservice ownership subjects.

The separate userservice signs JWTs. Other services verify those JWTs using the userservice public key/JWKS. This avoids making the execution backend responsible for password storage and login availability. If the serverless backend is under heavy invocation load, users should still be able to log in, refresh tokens, and manage accounts.

15.7 Direct runner upload tradeoff

Direct runner upload was introduced because Docker copy/export showed up in benchmarks. The idea is simple: if the function container already has the output files, it can validate and upload declared files directly instead of making the worker ask Docker to package those files and move them out.

The tradeoff is security and trust boundary. A runner with network access must receive scoped credentials and must not be allowed to upload arbitrary files for arbitrary invocations. The current design limits that risk by tying uploads to invocation ID, dispatch attempt, completion ID, declared output names, size limits, and backend-side verification before commit.

15.8 Warm containers before resident runners

The first warm-container design reused a container, but it still paid too much per-invocation setup inside the container. The bigger improvement came from the resident runner model. In that model the runtime process can remain ready to execute repeated requests for the

same function version, so tiny functions no longer pay full Python startup and import cost on every call.

The benchmark results reflect this. No-op functions benefited dramatically because most of their previous latency was platform overhead. Three-second sleep functions improved less because user code already occupied a large part of the total time.

16 Current Bottlenecks and Improvement Plan

16.1 Most important bottlenecks

The current bottlenecks are not all in one place. The biggest candidates are:

1. **Docker daemon contention under concurrent invocation.** The worker-concurrency sweep showed that raising concurrency from 8 to 12 or 16 did not produce a clean throughput win. The likely cause is not the orchestrator. Claim and completion were measured in tens of milliseconds. The heavier pressure is local to the worker host: Docker create/start, image checks, filesystem work, cleanup, and container lifecycle operations.
2. **Image pull and image check latency.** Workers pull from a single registry on the control plane. If a worker does not have the function image locally, invocation startup can pay a pull or manifest-check penalty. This is why image-cache-aware placement is a strong future scheduling improvement.
3. **Cold container create/start.** Cold paths still need Docker to allocate container metadata, configure mounts, attach the image filesystem, prepare networking mode, and start the process. This is not the same as user code execution.
4. **Python startup and imports.** Dependency-heavy Python functions can spend real time importing modules. Resident warm runners reduce this by keeping the interpreter and imported module state alive for repeated calls to the same function version.
5. **Output artifact handling.** Functions with declared outputs may need validation, upload, staged commit, and ZIP inclusion. Skip-export and direct runner upload reduce this, but output-producing functions still do more work than JSON-only functions.
6. **Warm-pool release and cleanup.** Benchmark timing showed warm-pool release and sandbox cleanup can become visible under burst load. This is not user code; it is the cost of preparing a reused container for the next invocation and cleaning invocation-specific state.
7. **Observability gaps on worker machines.** Metrics are present, but worker-side log shipping is not as complete as the control plane. That makes diagnosing distributed worker failures harder than it should be.

8. **Missing production-grade resource controls.** The prototype has timeout and artifact limits, but stronger CPU, process-count, memory, and disk controls are still needed before this is safe as a multi-tenant platform.

Bottleneck	Evidence	First improvement
Docker lifecycle pressure	Worker concurrency 12/16 did not improve throughput cleanly	Keep worker concurrency near measured saturation; reduce per-invocation Docker operations
Image cache misses	Pull/check phase appears in worker timing	Add image-cache-aware scheduling and proactive image pull after successful builds
Python import cost	Runner module import reached p95 values above one second in mixed runs	Use resident warm runners for eligible functions
Output handling	Output upload/export appears in p95 even when median is zero	Skip export when no declared outputs; improve direct runner upload
Warm-pool release	Release p95 exceeded one second in mixed-320 c12	Reduce cleanup scope and separate pool bookkeeping from slow Docker calls
Polling delay	Terminal latency can exceed worker duration	Add clearer status read semantics and optionally event-based client updates later

16.2 Highest-value improvements

The next improvements should be selected by expected user impact:

- Make the resident warm path the normal fast path for eligible functions.
- Add image cache-aware placement so invocations prefer workers that already have the image.
- Continue skipping all output artifact work for functions with no declared outputs.
- Improve direct upload for output-producing functions and measure it separately.
- Add worker-side log shipping to Loki.
- Add CPU and process-count limits before treating the system as multi-tenant.
- Build a dependency strategy for compiled wheels and system packages.

16.3 Research direction

If the project aims to become closer to commercial serverless performance, Docker alone may become a limitation. The next research comparison would be: optimized Docker warm containers, a resident runtime model, Firecracker microVMs, and language-specific worker processes. The prototype should not jump to a heavier isolation technology until the Docker warm path is fully measured, because the current bottlenecks include several removable prototype costs.

16.4 Dependency and wheel strategy

Python dependencies are a serious future risk. Packages such as NumPy often work because prebuilt wheels exist for the target Python version and Linux architecture. Other packages require system libraries, compilers, headers, or native build tools. The current platform lets a user provide `requirements.txt`, but it does not yet give them a structured way to request operating-system packages.

There are several possible fixes:

- **Curated runtimes:** provide prebuilt base images such as Python + NumPy/Pandas, Python + Pillow, or Python + PDF tools. This is predictable and fast but less flexible.
- **Admin-approved system packages:** allow selected apt packages in function build configuration. This is flexible but requires security review and cache strategy.
- **Build templates:** define named build profiles that install known native dependencies. This is a good middle ground for a prototype.
- **Full custom Dockerfile:** most flexible, but it makes the platform closer to a container hosting service and increases security risk.

For the current prototype, curated runtimes and build templates are the safest next step. They solve many real dependency problems without giving every user arbitrary operating-system build control.

17 Reading the Repository

17.1 Recommended reading order

A new developer should read the repository in this order:

1. `PROJECT_GUIDE.md`
2. `PROJECT_STATUS_REPORT.md`
3. `README.md`
4. `docs/frontend_status_results_api.md`
5. `docs/orchestrator_migration_contract.md`
6. `docs/first_distributed_deployment.md`
7. `worker/executor.py`
8. `worker/worker.py`
9. `worker/warm_pool.py`
10. `scheduler/production_orchestrator.py`
11. `backend/apps/functions/`
12. `backend/apps/invocations/`
13. `backend/apps/jobs/`
14. `userservice/apps/accounts/`
15. `benchmark-results/`

17.2 Backend code

The backend is a Django project. The most important apps are:

- apps/functions** Function models, source upload, build status, invocation tokens, active versions, image ledger.
- apps/invocations** Invocation records, input files, output files, log artifacts, staged completions, response modes, synchronous invocation.

apps/jobs	Durable job rows, outbox, projection, V1 retirement, retention, reconciliation.
apps/accounts	Compatibility bridge for userservice JWT verification and local shadow users.
apps/workers	Worker model and admin-only worker API.

17.3 Worker code

The worker directory contains:

worker.py	Worker process, heartbeat loop, queue consumption, thread pool.
executor.py	Docker invocation executor, sandbox handling, result collection, timings.
builder.py	Build context preparation, Dockerfile/runner generation, image build and push.
warm_pool.py	Warm container lifecycle, resident runner reuse, eviction policy.
backend_client.py	Protected backend calls for payloads, files, outputs, and reports.
orchestrator_client.py	Claim and completion calls to the orchestrator.

17.4 Scheduler and orchestrator code

The scheduler directory contains the older scheduler and the V2 production orchestrator. The V2 direction is centered around:

production_orchestrator.py	Dispatch, claim, completion, recovery, and metrics endpoints.
orchestrator_state.py	Redis-backed state transitions and helper logic.
orchestrator_workers.py	Worker heartbeat and worker selection support.
invocation_finalizer.py	Commit staged artifacts and complete finalization.
orphan_image_cleaner.py	Clean unused images from the registry.

17.5 Frontend code

The frontend is a React/Vite application. It presents a landing page, login and signup, a gamified step-by-step function creation flow, source editor, build status, invocation controls, function-token management, and user-facing tutorial pages. The frontend should rely on the documented links and response modes instead of internal V1/V2 coordination details.

18 Conclusion

The platform has moved from a simple proof of concept into a working distributed prototype. It supports the main serverless lifecycle: account login, function creation, source upload, image build, asynchronous and synchronous invocation, input files, output files, logs, ZIP download, function tokens, retention cleanup, observability, and distributed workers.

The strongest engineering part of the project is the gradual orchestration design. The system did not jump directly to a complicated distributed architecture. It first established the basic path, then added durable jobs, worker-specific queues, heartbeats, recovery, backoff, dead letters, V2 Redis Streams, staged artifacts, finalization, projection, warm routing, resident warm containers, and distributed deployment. This is the right kind of path for a bachelor project because it shows implementation, measurement, correction, and iteration.

The main remaining work is also clear. The prototype needs stronger sandbox resource controls, better dependency/build handling, worker log shipping, further performance tuning, and production-grade hardening. The current bottleneck is no longer simply “the platform does not work.” The platform works. The next challenge is making it faster, safer, easier to operate, and more predictable under mixed real workloads.

A Selected API Reference

Endpoint	Purpose
POST /api/auth/register/	Register a userservice account.
POST /api/auth/token/	Log in and receive JWT tokens.
POST /api/auth/token/refresh/	Refresh access token.
GET /api/auth/me/	Read current account profile.
POST /api/functions/	Create function metadata.
GET /api/functions/	List owned functions.
GET /api/functions/{id}/source/	Read editable handler and requirements for the editor.
POST /api/functions/{id}/source/	Upload source bundle and queue build.
GET /api/functions/{id}/build-status/	Read frontend-friendly build state.
GET /api/functions/{id}/tokens/	List invocation tokens.
POST /api/functions/{id}/tokens/	Create invocation token.
POST /api/functions/{id}/tokens/{token_id}/rotate/	Rotate invocation token secret.
POST /api/functions/{id}/tokens/{token_id}/revoke/	Revoke invocation token.
POST /api/functions/{id}/invoke/	Create asynchronous invocation.
POST /api/functions/{id}/invoke-sync/	Run synchronous invocation when eligible.
GET /api/invocations/{id}/	Poll invocation status and result.
GET /api/invocations/{id}/outputs/	List output artifacts.
GET /api/invocations/{id}/outputs/{file_id}/download/	Download one output artifact.
GET /api/invocations/{id}/download/	Download invocation ZIP.

B Example Function With Output File

Listing B.1: Function that returns JSON and writes one declared output.

```
1 from pathlib import Path
2
3 def main(event, context):
4     numbers = event.get("numbers", [])
5     total = sum(numbers)
6     name = event.get("name", "anonymous")
7
8     output_dir = Path(context["output_dir"])
9     output_dir.mkdir(parents=True, exist_ok=True)
10    (output_dir / "report.txt").write_text(
11        f"Processed {name} with total={total}",
12        encoding="utf-8",
13    )
14
15    print(f"processed request for {name}")
16
17    return {
18        "sum": total,
19        "message": f"Hello, {name}!",
20    }
```

C Glossary

Artifact	A file produced or stored by the platform, such as a source bundle, input file, output file, log, or ZIP.
Build	The process of converting source into a Docker image.
Cold start	An invocation that must create/start a new container or load a runtime from scratch.
Dispatch attempt	A fenced attempt number assigned by the orchestrator.
Finalizer	Service that commits staged artifacts after completion is accepted.
Function token	Function-scoped invocation secret, separate from JWT.
Invocation	One execution request for a built function.
Orchestrator	Service responsible for active V2 job coordination.
Projection	Copying authoritative Redis events into PostgreSQL read models.
Staged artifact	Uploaded artifact that physically exists but is hidden until final commit.
Warm container	Reusable container already prepared for a specific function version.